

# IO ECO 7408: Problem Set 1 – Binary Logit Demand Estimation for Ride-Sharing Services

Instructor: Prof. Douglas Turner

Student: C.H.

```
In [25]: globals().clear()
```

```
In [26]: import sys
import numpy as np
import pandas as pd
import statsmodels.api as sm
from statsmodels.discrete.discrete_model import Logit
from scipy.optimize import minimize

import matplotlib.pyplot as plt
```

## Questions

### 1. Data Setup and Variable Construction

(a) Load the dataset into Python using pandas. Print the first few rows to inspect the data.

```
In [27]: df = pd.read_csv('/Users/terrylu/Desktop/UF/Courses/2025-2026/IO/IO_Code/IO_Python/IO_Assignment1.html')
print(df.head(5))
```

	wait_time_uber	wait_time_lyft	price_uber	price_lyft	age	choice_
is_uber						
0	18.528105	15.619448	8.296000	26.216232	29	
1						
1	15.800314	13.525088	15.646098	16.062803	38	
0						
2	16.957476	11.926160	18.613206	21.884515	53	
1						
3	19.481786	13.875490	23.575972	20.009804	50	
0						
4	18.735116	11.800978	9.246689	24.669068	21	
1						

(b) Construct the following new variables:

- wait\_diff: Uber wait time minus Lyft wait time
- price\_diff: Uber price minus Lyft price

Also include choice\_is\_uber in your dataframe.

Hint: A more positive value of wait\_diff or price\_diff implies Uber is worse than lyft in that attribute.

```
In [28]: df['wait_diff'] = df['wait_time_uber'] - df['wait_time_lyft']
df['price_diff'] = df['price_uber'] - df['price_lyft']
```

## 2. Estimating the Logit Model by Hand

(a) Using only numpy and scipy.optimize.minimize, estimate a binary logit model of the form:

Write out the log-likelihood function manually and optimize it using BFGS.

```
In [29]: img = plt.imread('/Users/terrylu/Desktop/UF/Courses/2025-2026/IO/IO_Code/IO_Python/IO_Assignment1.html')
plt.imshow(img); plt.axis("off");
```

$$\Pr(\text{choose Uber}) = \frac{\exp(\beta_1 \cdot \text{wait\_diff} + \beta_2 \cdot \text{price\_diff})}{1 + \exp(\beta_1 \cdot \text{wait\_diff} + \beta_2 \cdot \text{price\_diff})}$$

Write out the log-likelihood function manually and optimize it using BFGS.

```
In [30]: y = df['choice_is_uber']
X = df[['wait_diff', 'price_diff']]

def log_likelihood(beta): # define the log-likelihood function
    V = X @ beta # systematic utility
    P = np.exp(V) / (1 + np.exp(V))
    # avoid log(0)
    esp = 1e-12
    P = np.clip(P, esp, 1-esp)
    # log-likelihood
    ll = y * np.log(P) + (1 - y) * np.log(1 - P)
```

```

    return -np.sum(ll) # negative log-likelihood for minimization

# initial guess for beta
beta_init = np.zeros(X.shape[1])

# optimize the log-likelihood function
result_1 = minimize(log_likelihood, beta_init, method = 'BFGS')

```

(b) Report the estimated coefficients. Do the signs of the estimates make sense? Why?

```

In [31]: # print the results
beta_hat_1 = result_1.x
print(f"Estimated coefficients (by hand):\n{beta_hat_1}")

```

```

Estimated coefficients (by hand):
[-0.5241069 -2.03895639]

```

Yes, those coefficients make sense.

Both of them are negative, which inflects that, longer waiting time and higher price will decrease the probability of uber being chosen.

(c) Why can't age be included in the above model?

Because of the endogeneity.

Of course, age can affect utility.

But the model also has an error term, which in fact captures many unobserved or omitted variables, for example:

1. how user-friendly the platform is for older people.
2. whether the platform supports credit cards (which may matter more for adults).

Since these factors are correlated with age, including age directly in the model would lead to endogeneity problems.

### 3. Estimation Using statsmodels

(a) Re-estimate the same binary logit model using statsmodels.Logit. Compare the results with your manual implementation. Hint: They should be similar if done correctly. Due to numerical differences, they may not be exactly the same.

```

In [32]: logit_model_1 = Logit(y, X)

```

```
logit_result_1 = logit_model_1.fit(dis=1)

print("\n\nEstimation Using statsmodels\n\n\n", logit_result_1.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.066951
      Iterations 12
```

Estimation Using statsmodels

```

                                Logit Regression Results
=====
Dep. Variable:                choice_is_uber    No. Observations:
5000
Model:                        Logit            Df Residuals:
4998
Method:                       MLE             Df Model:
1
Date:                        Mon, 22 Sep 2025    Pseudo R-squ.:
0.9033
Time:                        20:47:43          Log-Likelihood:
-334.75
converged:                    True             LL-Null:
-3463.0
Covariance Type:              nonrobust         LLR p-value:
0.000
=====
=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
wait_diff      -0.5241      0.047     -11.167      0.000     -0.616
-0.432
price_diff     -2.0390      0.117     -17.451      0.000     -2.268
-1.810
=====
=====
```

Possibly complete quasi-separation: A fraction 0.66 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

In [33]: *### With Robust Covariance Matrix*

```
logit_result_2_robust = logit_model_1.fit(dis=1, cov_type='HC0')

print("\n\nEstimation Using Statsmodels With Robust Covariance Matrix
```

Optimization terminated successfully.  
 Current function value: 0.066951  
 Iterations 12

### Estimation Using Statsmodels With Robust Covariance Matrix

```

                                Logit Regression Results
=====
Dep. Variable:                choice_is_uber    No. Observations:
5000
Model:                        Logit            Df Residuals:
4998
Method:                       MLE             Df Model:
1
Date:                        Mon, 22 Sep 2025    Pseudo R-squ.:
0.9033
Time:                        20:47:43          Log-Likelihood:
-334.75
converged:                   True              LL-Null:
-3463.0
Covariance Type:             HC0              LLR p-value:
0.000
=====
=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
wait_diff      -0.5241      0.049     -10.595      0.000     -0.621
-0.427
price_diff     -2.0390      0.122     -16.671      0.000     -2.279
-1.799
=====
=====

```

Possibly complete quasi-separation: A fraction 0.66 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

Almost the same.

(b) What is the implied willingness to wait one fewer minute (in dollars)?

Hint: Think about how much a consumer is willing to pay to save a minute of wait time.

"Willingness to wait one fewer minute (in dollars) means": keeping the overall

utility unchanged, if time increases by one minute, then the price must decrease by how much?

```
In [34]: img = plt.imread('/Users/terrylu/Desktop/UF/Courses/2025-2026/IO/IO_Code/IO_Python/IO_Assignment1.html')
plt.imshow(img); plt.axis("off");
```

$$U_{\text{bar}} = \beta_{\text{wait}} \cdot \text{time} + \beta_{\text{price}} \cdot \text{price}_1$$

$$U_{\text{bar}} = \beta_{\text{wait}} \cdot (\text{time} + 1) + \beta_{\text{price}} \cdot \text{price}_2$$

Difference:

$$0 = \beta_{\text{wait}} \cdot 1 + \beta_{\text{price}} \cdot \Delta \text{price}$$

$$\Delta p = -\frac{\beta_{\text{wait}}}{\beta_{\text{price}}}$$

```
In [35]: beta_wait = logit_result_2_robust.params['wait_diff']
beta_price = logit_result_2_robust.params['price_diff']

delta_price = - (beta_wait / beta_price)
print(f"The implied willingness to wait one fewer minute (in dollars)
```

The implied willingness to wait one fewer minute (in dollars) is: -0.2570

## 4. Re-estimate with a Product Fixed Effect

```
In [36]: img = plt.imread('/Users/terrylu/Desktop/UF/Courses/2025-2026/IO/IO_Code/IO_Python/IO_Assignment1.html')
plt.figure(figsize=(10, 10))
plt.imshow(img); plt.axis("off");
```

## 1. Utility framework

We have two alternatives, Uber and Lyft:

$$U_{Uber} = \alpha_{Uber} + \beta_1 \cdot wait_{Uber} + \beta_2 \cdot price_{Uber}$$

$$U_{Lyft} = \alpha_{Lyft} + \beta_1 \cdot wait_{Lyft} + \beta_2 \cdot price_{Lyft}$$

## 2. Difference in utilities

For binary logit we only need:

$$\begin{aligned} V &= U_{Uber} - U_{Lyft} \\ &= (\alpha_{Uber} - \alpha_{Lyft}) + \beta_1 (wait_{Uber} - wait_{Lyft}) + \beta_2 (price_{Uber} - price_{Lyft}) \end{aligned}$$

Define:

$$\beta_0 = \alpha_{Uber} - \alpha_{Lyft}$$

So:

$$V = \beta_0 + \beta_1 \cdot wait\_diff + \beta_2 \cdot price\_diff$$

Here  $\beta_0$  is simply the constant capturing Uber's inherent preference relative to Lyft.

## 3. Dummy-variable representation

Some textbooks prefer to write instead:

$$V = \gamma \cdot D_{Uber} + \beta_1 \cdot wait\_diff + \beta_2 \cdot price\_diff$$

- $D_{Uber} = 1$  if the alternative is Uber, 0 if Lyft.
- Since we are already working with the **difference Uber – Lyft**, the relevant case is always Uber, so  $D_{Uber} = 1$ .

Thus:

$$V = \gamma + \beta_1 \cdot wait\_diff + \beta_2 \cdot price\_diff$$

(a) Re-estimate the logit model, this time including a dummy variable which is 1 if the product is uber and 0 for lyft.

```
In [37]: X_inter = sm.add_constant(df[['wait_diff', 'price_diff']])
```

```

logit_model_2 = Logit(y, X_inter)

logit_result_2 = logit_model_2.fit(dis=1)

print("\n\nEstimation with Interception\n\n\n", logit_result_2.summary(

```

Optimization terminated successfully.  
 Current function value: 0.066753  
 Iterations 12

Estimation with Interception

```

                                Logit Regression Results
=====
Dep. Variable:                choice_is_uber    No. Observations:
5000
Model:                        Logit            Df Residuals:
4997
Method:                       MLE             Df Model:
2
Date:                         Mon, 22 Sep 2025   Pseudo R-squ.:
0.9036
Time:                         20:47:44         Log-Likelihood:
-333.77
converged:                    True             LL-Null:
-3463.0
Covariance Type:              nonrobust         LLR p-value:
0.000
=====
=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
const                0.1384      0.099      1.401      0.161      -0.055
0.332
wait_diff            -0.5251      0.047     -11.151      0.000      -0.617
-0.433
price_diff           -2.0450      0.117     -17.433      0.000      -2.275
-1.815
=====
=====

```

Possibly complete quasi-separation: A fraction 0.66 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

(b) Why might it be important to include this dummy in the model? What does the



intercept capture?

Without the intercept, the model forces the probability of choosing Uber to be 0.5 when  $\text{wait\_diff}=0$  and  $\text{price\_diff}=0$ .

With the intercept, the model allows for this baseline preference to differ from 0.5.

For example, the intercept here is 0.1384, which means that when  $\text{wait\_diff}=0$  and  $\text{price\_diff}=0$ , the probability of choosing Uber is 0.5345, which is greater than 0.5. Which indicates that, Uber has a slight baseline preference advantage over Lyft.

## 5. Prediction

Suppose an individual faces the following options:

- Uber: 22-minute wait, \$7
- Lyft: 14-minute wait, \$10

(a) Construct the appropriate values for  $\text{wait\_diff}$  and  $\text{price\_diff}$ .

```
In [38]: wait_diff_pre = 22 - 14
price_diff_pre = 7 - 10
```

(b) Use your most recent model (with the intercept) to predict the probability that the individual chooses Lyft.

```
In [39]: x_pre = [[1, wait_diff_pre, price_diff_pre]]
y_pre = logit_result_2.predict(x_pre)
y_pre_lyft = 1 - y_pre
print(f"predicted probability of choosing Uber:\n {y_pre}")
print(f"predicted probability of choosing Lyft:\n {y_pre_lyft}")
```

```
predicted probability of choosing Uber:
[0.88823384]
predicted probability of choosing Lyft:
[0.11176616]
```

## 6. Heterogeneity by Age

(a) Add a new interaction term  $\text{age\_wait\_diff}$  equal to  $\text{wait\_diff} * \text{age}$ .

```
In [40]: df['age_wait_diff'] = df['wait_diff'] * df['age']
```

(b) Re-estimate the model using wait\_diff, price\_diff, and age\_wait\_diff (plus a constant).

```
In [41]: x_6 = df[['wait_diff', 'price_diff', 'age_wait_diff']]
x_6 = sm.add_constant(x_6)

logit_model_6 = Logit(y, x_6)
logit_result_6 = logit_model_6.fit(dis=0)
print("Estimation with Age & Interaction\n\n\n", logit_result_6.summary)
```

## Estimation with Age &amp; Interaction

```

                                Logit Regression Results
=====
=====
Dep. Variable:                choice_is_uber    No. Observations:
5000
Model:                        Logit            Df Residuals:
4996
Method:                       MLE             Df Model:
3
Date:                         Mon, 22 Sep 2025   Pseudo R-squ.:
0.9036
Time:                         20:47:52         Log-Likelihood:
-333.76
converged:                    True             LL-Null:
-3463.0
Covariance Type:              nonrobust         LLR p-value:
0.000
=====
=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
const                0.1386        0.099        1.402      0.161      -0.055
0.332
wait_diff            -0.5323        0.115       -4.649      0.000      -0.757
-0.308
price_diff           -2.0449        0.117      -17.431      0.000      -2.275
-1.815
age_wait_diff         0.0002        0.003        0.070      0.944      -0.005
0.005
=====
=====

```

Possibly complete quasi-separation: A fraction 0.66 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

The coefficient on age\_wait\_diff is close to zero and statistically insignificant.

This implies that in this dataset, the effect of wait time on the probability of choosing Uber does not vary systematically with age.

## 7. Reflection

(a) Suppose a third option was added to the dataset. Specifically, suppose consumers also had the option of choosing uber black (a premium version of

uber). Briefly explain one limitation of the logit model in this context.

One limitation of the logit model is the IIA (Independence of Irrelevant Alternatives) assumption: the relative choice probability between any two options does not change when a third option is added or removed.

In other words, in the logit model, the choice share between Uber and Lyft would remain the same regardless of whether Uber Black exists.

However, In reality, this assumption clearly does not hold, because Uber Black is much closer a substitute to Uber than to Lyft (they are available within the same app). As a result, many consumers who originally chose Uber would likely switch to Uber Black, rather than randomly switching away from both Uber and Lyft. However, the logit model would incorrectly predict that the introduction of Uber Black takes market share proportionally from both Uber and Lyft, which is unrealistic.